
Data Structures and Algorithms

Lecture 10: Indexing & Advanced Trees

Outline

■ Indexing

- Linear index
- Tree-based index
 - B-trees, B⁺-trees

■ Advanced Trees

- Tries
- Balanced trees
 - AVL tree, Red-black tree, BB(α) tree, Splay tree
- Spatial data structures
 - K-D tree, PR quadtree

Application limitations of Hash

- Hash provides excellent performance for insert, search, and delete, i.e.,
 - Time complexity $\Theta(1)$ on average
- But hash has some application limitations:
 1. Do not support duplicate keys
 2. Only provide exact-search, but not range search
 - E.g., search the students with their height between 1.7m and 1.75m
 3. Do not support efficient searching the record with the minimum or maximum key


What is Index ?

- Index provides following operations:
 - efficient **Insert** (with duplicate keys): $\Theta(\log n)$
 - efficient **exact-search**: $\Theta(\log n)$
 - efficient **range-search**, time is related to the range, but usually is **much shorter** than $\Theta(n)$
 - Efficient **minimum / maximum search**: $\Theta(\log n)$
 - Efficient delete: $\Theta(\log n)$
- Index is designed for a **large collection** of records stored on **disks**, where the disk access time is **much slower** than memory access time.

Linear indexing

- A **linear index** is an index file organized as a sequence of key-value pairs where the keys are in sorted order and the pointers either
 - ❑ (1) point to the position of the complete record on disk,
 - ❑ (2) point to the position of the primary key in the primary index,
 - ❑ or (3) are actually the value of the primary key.

Linear Index




Database Records

- If the database contains enough records, the linear index might be **too large** to store in main memory. -> **expensive!**

Index techniques has many similarities with BST

- A **binary search tree** (BST) is a special binary tree, iff
 - For each node, assume the node value is K
 - The values of the nodes in its **left subtree** are $< K$
 - The values of the nodes in its **right subtree** are $\geq K$



Index techniques

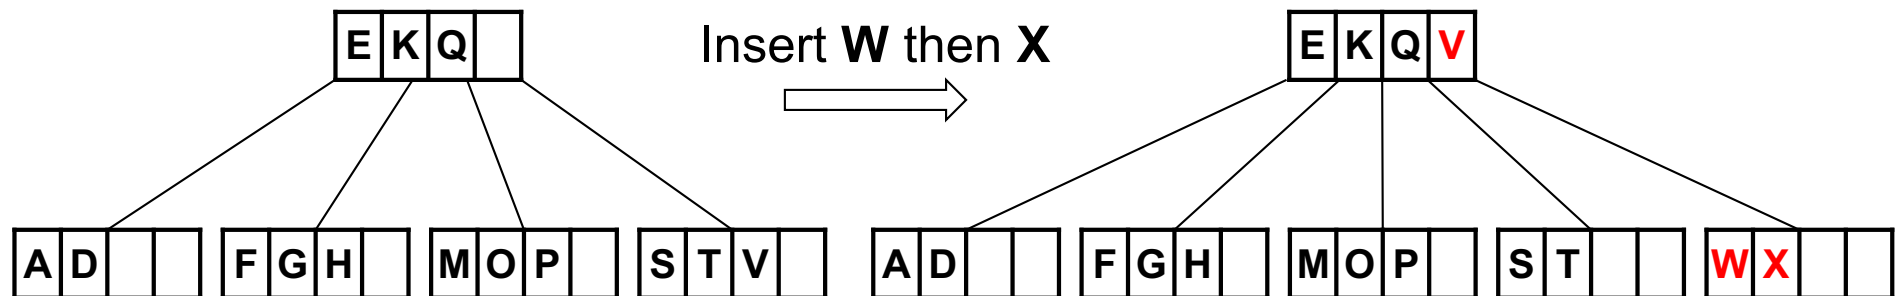
- Two Tree-based indexing techniques:
 - B trees
 - B+ trees
- Why not adopt a binary search tree (BST) for index?
 - A BST may **not** be **balanced**
 - E.g., One **subtree** has **many** nodes, while the other has **a few** nodes, **poor performance**
 - The **depth** of a **balanced** BST is **still large**
 - Need about $\log_2 n$ searches, and possible $\log_2 n$ times of **disk accesses**, while a disk access is very **time-consuming**. This is unacceptable.

B tree

- B tree is a height **Balanced** tree.
- A B tree of **order m** is defined to have the following properties:
 - The **root** is has **at least two children** or either a **leaf**.
 - Each **node**, except for the root and the leaves, has between **$\lceil m/2 \rceil$** and **m children**.
 - Typically, **m** will be fairly **large**, e.g., **$m=100$**
 - **All leaves** are at the **same level** in the tree, so the tree is always **height balanced**.
 - The data values in each node are in ascending order.

Node insertion in a B tree

- Insertion follows similar logic to the BST.
 - ▣ **Basic idea:** search for the appropriate leaf, add the new value, then split and promote as necessary.



1. Insert **W** (just fills the leaf),
2. Then insert **X** (would cause the right-most leaf to split, and **V** to be promoted to the root),

Node deletion in a B tree


- The process of node deletion is **similar** to that in BST
- Deletion from a **Leaf node**
 - May drop the number of data values in the node below the mandatory floor ("underflow").
 - In this case, the leaf must borrow a value from an adjacent sibling node if node have a value to spare, or be merged with an adjacent sibling node.
- Deletion from an **Internal node**
 - Accomplished by reducing it to the former case.

Search in a B tree

- Main steps:
 1. Perform a binary search on the records in the current node.
 - If a record with the search key is found, then return that record.
 - If the current node is a leaf node and the key is not found, then report an unsuccessful search.
 2. Otherwise, follow the proper branch and repeat the process.

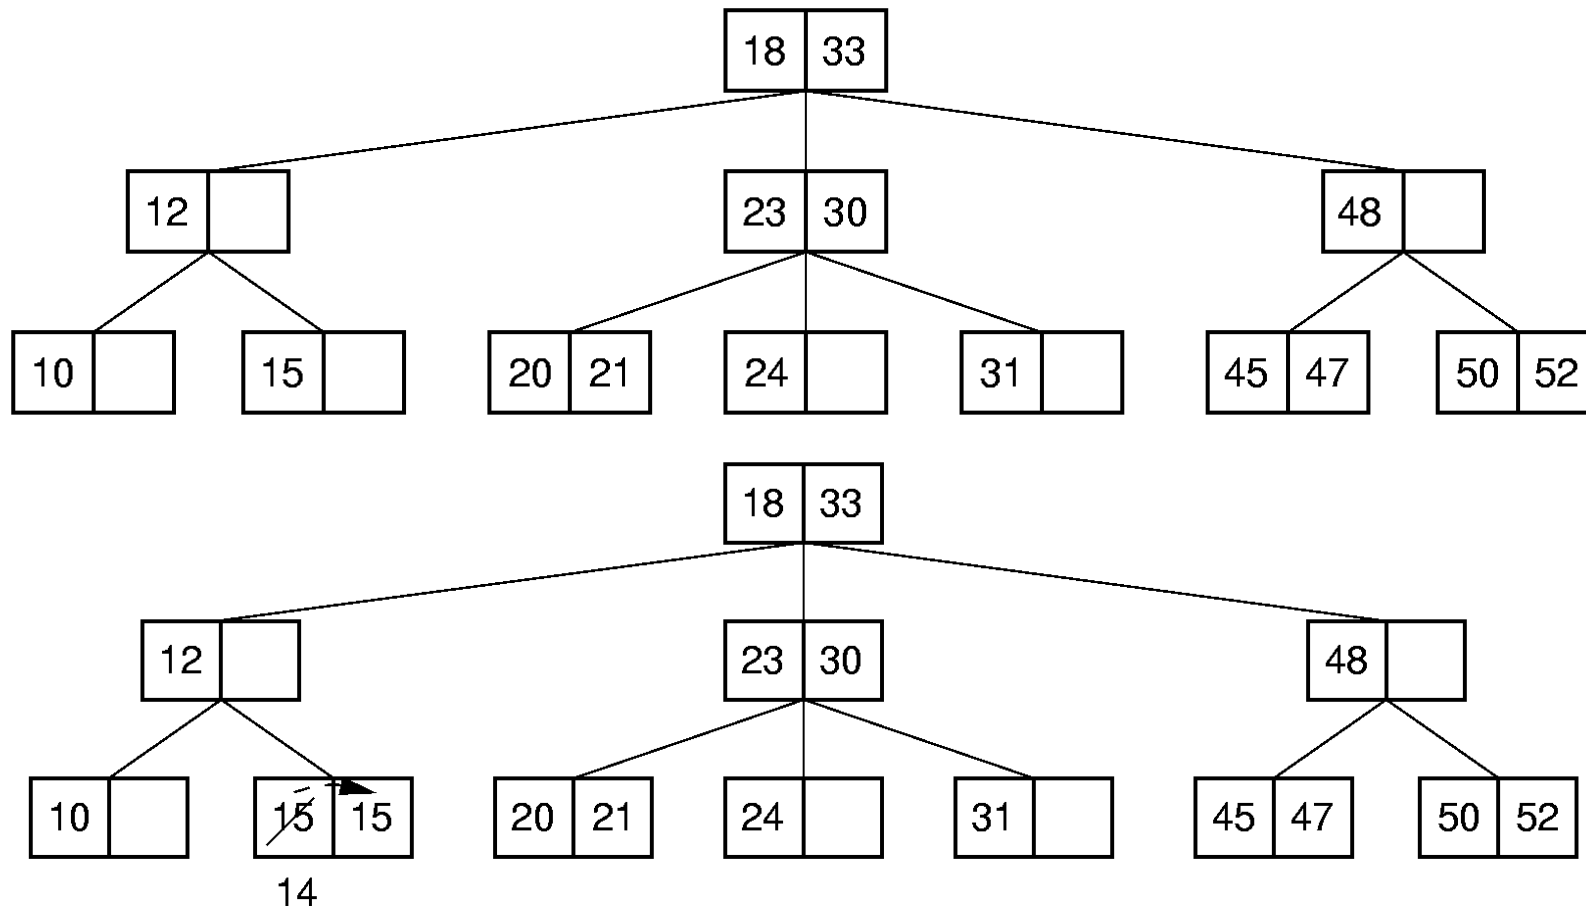
B tree example: 2-3 tree, i.e., $m = 3$

- Each internal nodes in a 2-3 tree has 2 or 3 children
 - A node contains **one or two keys**
 - All leaves are at the same level
- The 2-3 Tree has a property analogous to the BST:
 - **left subtree** $< 1^{\text{st}}$ key;
 - 1^{st} key \leq **mid subtree** $< 2^{\text{nd}}$ key;
 - **right subtree** $\geq 2^{\text{nd}}$ key



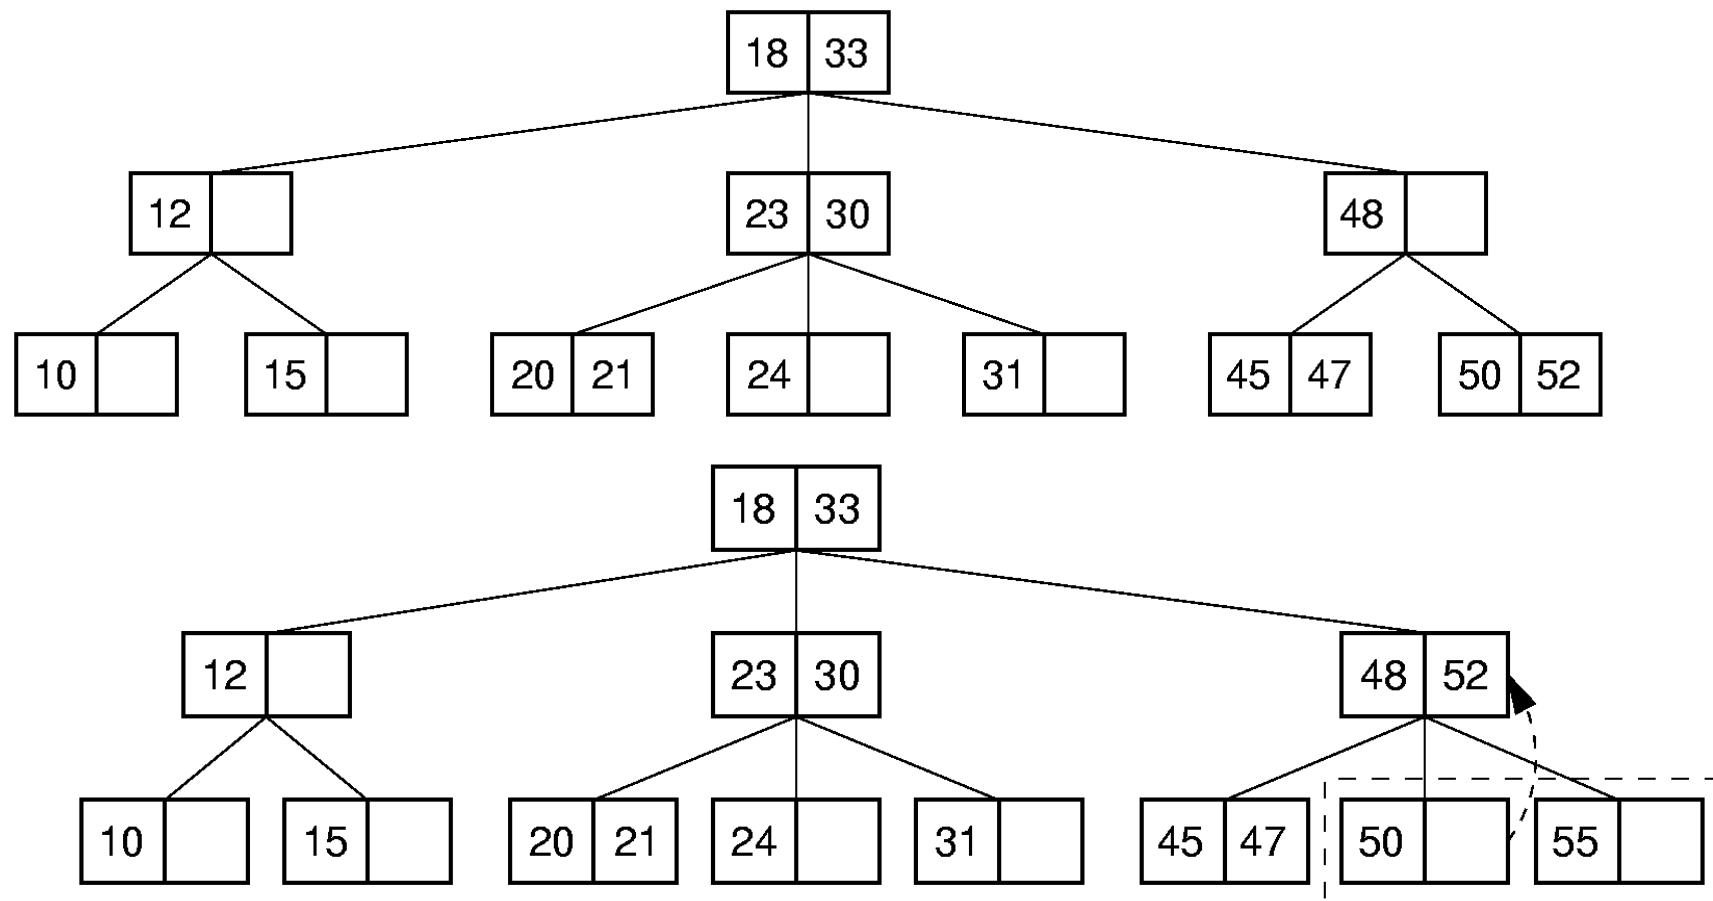
2-3 Tree

- The advantage of the 2-3 Tree over the BST is that it can be updated **at low cost**, e.g., insert **14**




2-3 Tree Insertion, insert 55


- Split the node has keys 50 and 52, and
- Promote the median of 50, 52, 55 to its parent



2-3 Tree Insertion, insert 19




- Split the node has 20,21, promote 20, to node has 23, 30
- Then split node has 23,30, promote 23 to root has 18, 33



2-3 Tree Insertion, insert 19

- Split the root has 18, 33 due to the insertion 23, and
- promote 23, by creating a new root
- The tree height increase by 1
- But all leaves at the same level



(c)

B⁺ Trees


- The most **commonly** implemented form of the B-Tree is the **B⁺ Tree**.
- B⁺ tree stores records **ONLY** at the leaf nodes.
- **Internal** nodes store **keys** to **guide** the **search**.
- **Leaf** nodes store actual records, or else keys and pointers to actual records.
- A **leaf** node can store no more than **m+1** records
- B+ tree supports **$\Theta(1)$** time to search the **previous** or **next** record, of a given record.

B⁺ Trees (cont'd)

- Define a B+ tree of **order m** as follows:
 - All data is stored at the leaf blocks
 - The root nodes is either:
 - A leaf block, or
 - An **m -way tree** with between **2** and **m** children
 - All other internal blocks are **m -way trees** with **$\lceil m/2 \rceil$** to **m** children
 - The internal blocks store up to $m - 1$ keys to guide the searching where key k denotes the smallest key in sub-tree k .


B⁺-Tree Example with order $m=4$

- Each internal node should have from $\lceil m/2 \rceil = 2$ to $m=4$ children
- A leaf has no more than $m+1=5$ records, but at least $\lceil (m+1)/2 \rceil = 3$ records
- Nodes in the same level are linked in order




B⁺-Tree Insertion

- Insert 55
- Similar the insertion in B tree





(a)



(b)

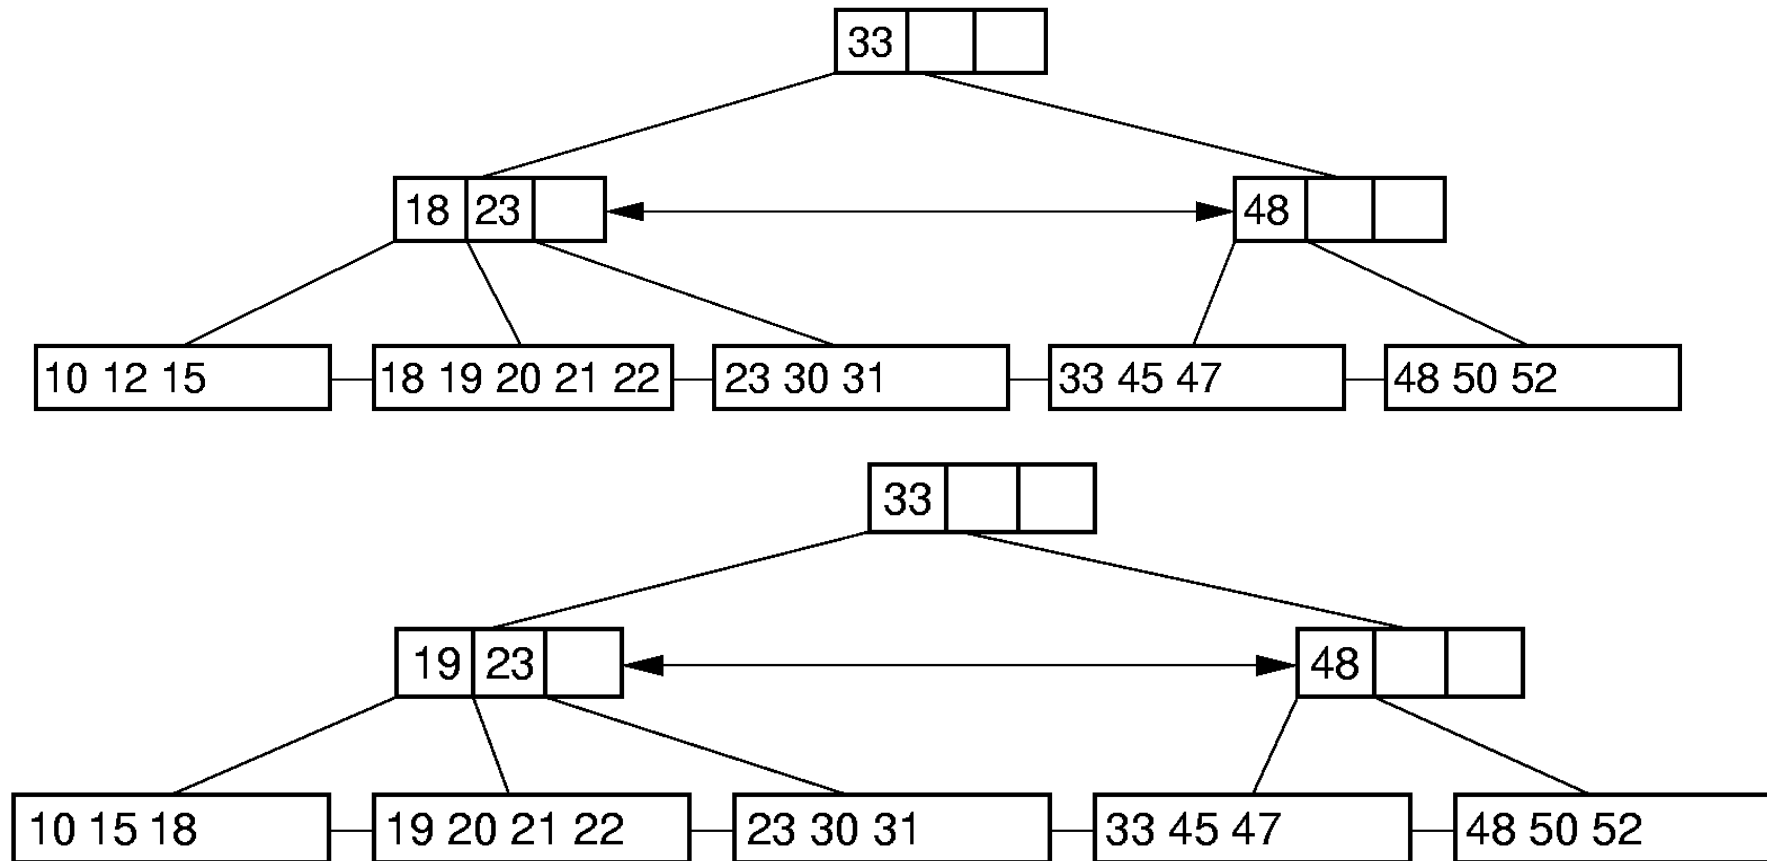
B⁺-Tree Deletion (1) - delete 18

- Just remove key 18 from its leaf node



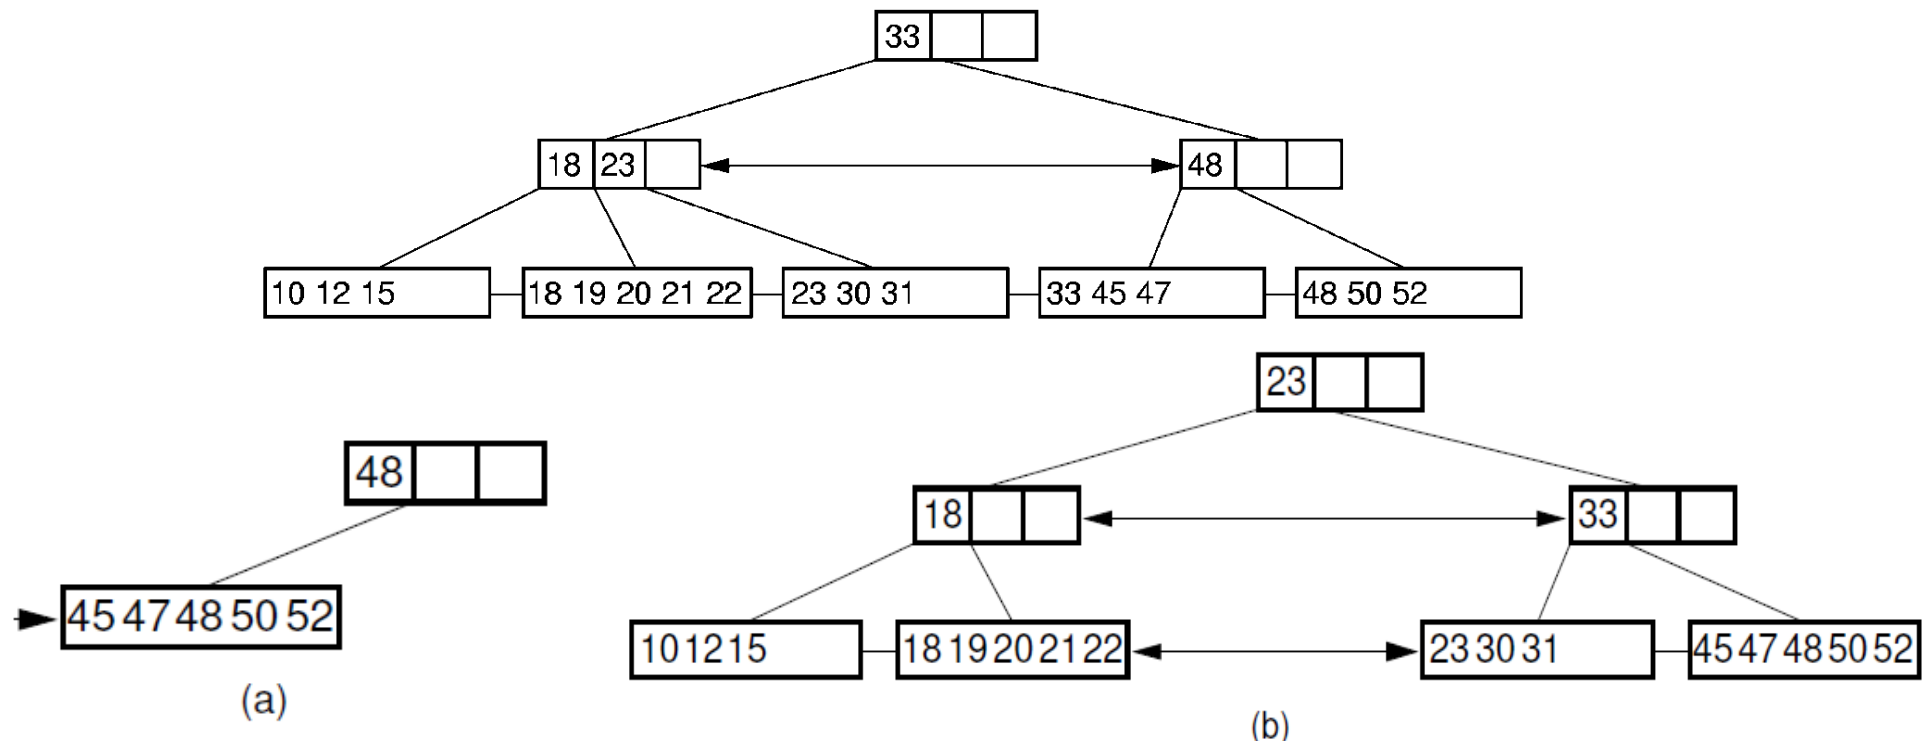
B⁺-Tree Deletion (2) - delete 12

- Borrow one node 18 from its sibling to make it **at least 3 nodes**



B⁺-Tree Deletion - delete 33

- Node having 33,45,47 cannot borrow from its siblings, merge with its one sibling node 48,50,52
- Node 48 has one less child, borrow one child from its sibling node having 18,23, modify guide keys



B* Trees

- A variant of B⁺ trees
- All nodes except the root are required to be at least 2/3 full rather than 1/2 full.
- Splitting transforms 2 nodes into 3, rather than 1 node into 2.
- Can be generalized to specify a fill factor of $(n+1)/(n+2)$; a Bⁿ tree.

B-Tree Analysis

- Asymptotic cost of search, insertion, and deletion of nodes from B-Trees is $\Theta(\log n)$.
 - Base of the log is the (average) branching factor of the tree.
- Example: Consider a B+-Tree of order 100 with leaf nodes containing $m=100$ records.
 - 1 level B+-tree: Min 0, Max 100
 - 2 level: Min: 2 leaves of 50 (100 records). Max: 100 leaves with 100 (10,000 records).
 - 3 level: Min 2 x 50 nodes of leaves, for 5000 records. Max: $100^3 = 1,000,000$ records.
 - 4 level: Min: 250,000 records ($2 * 50 * 50 * 50$). Max: $100^4 = 100$ million records.

Advanced Tree Structures

Advanced Tree Structures

1. Tries

2. Balanced trees

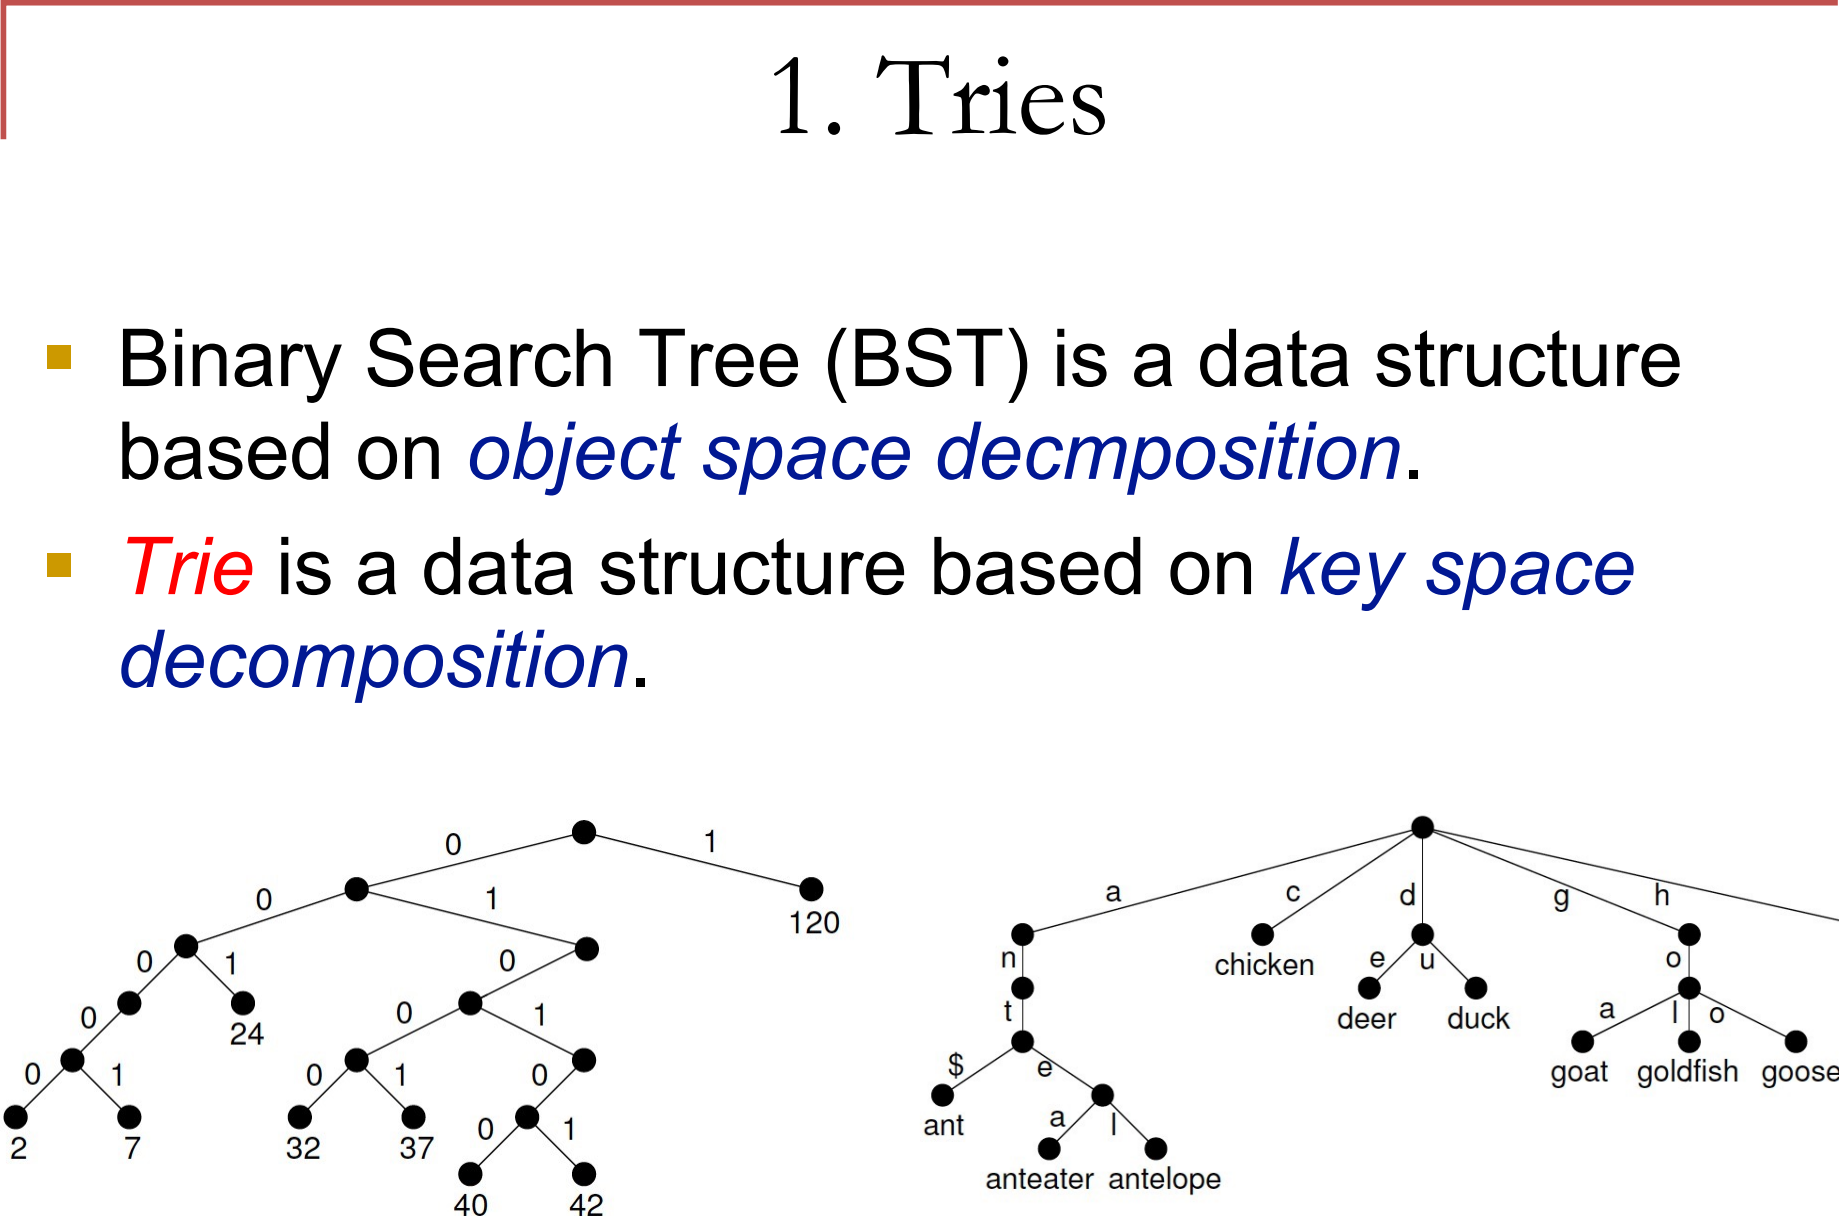
- ▣ AVL tree, Red-black tree, $BB(\alpha)$ tree, Splay tree

3. Spatial data structures

- ▣ K-D tree
- ▣ PR quadtree

1. Tries

- # 1. Tries
- Binary Search Tree (BST) is a data structure based on *object space decomposition*.
 - *Trie* is a data structure based on *key space decomposition*.



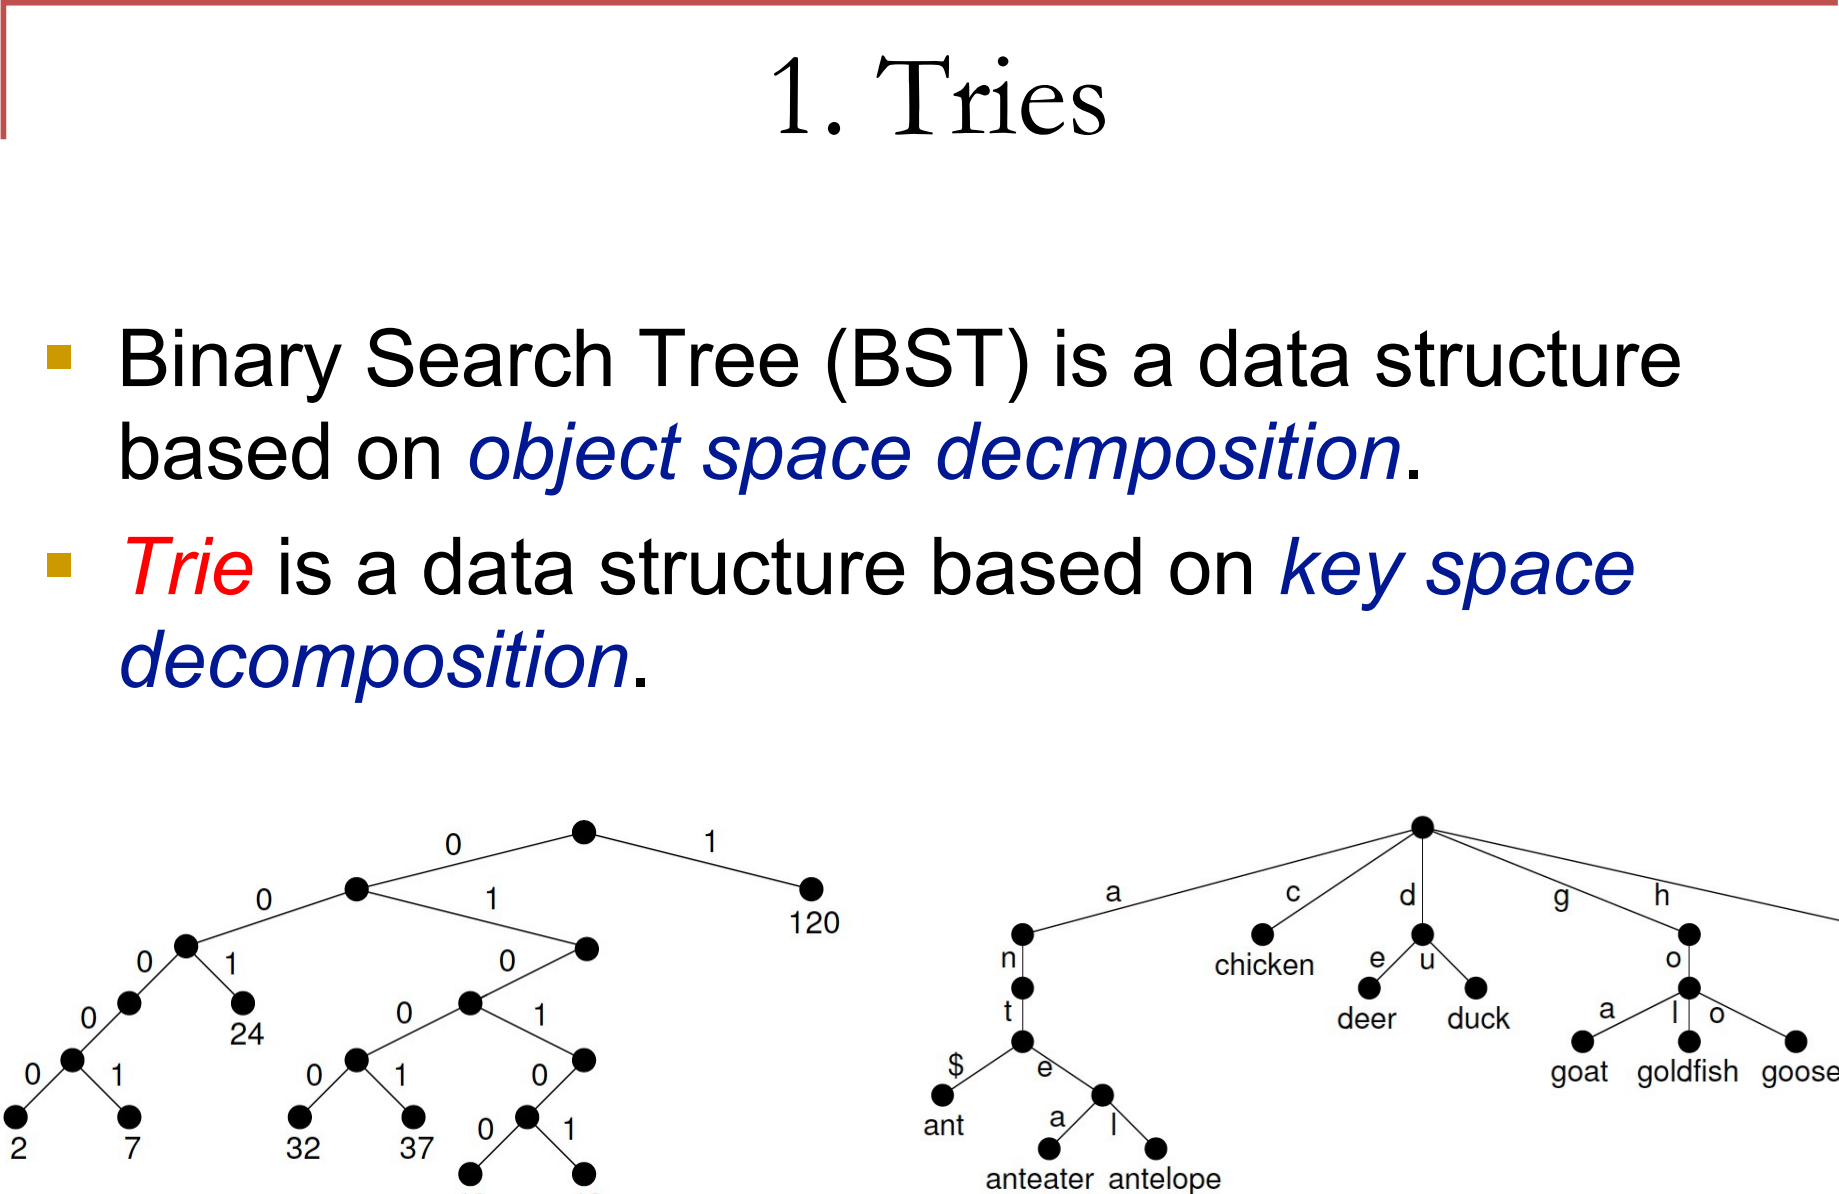
1. Tries

- Binary Search Tree (BST) is a data structure based on *object space decomposition*.
- *Trie* is a data structure based on *key space decomposition*.

(a) Binary trie

(b) alphabet trie

The diagram illustrates two types of trie structures. (a) Binary trie: A binary tree where internal nodes are black dots. Edges are labeled 0 and 1. Leaf nodes are labeled with integers: 2, 7, 24, 32, 37, 40, 42, and 120. (b) Alphabet trie: A tree where internal nodes are black dots. Edges are labeled with characters. Leaf nodes are labeled with words: ant, anteater, antelope, chicken, deer, duck, goat, goldfish, and goose.



1. Tries

- Binary Search Tree (BST) is a data structure based on *object space decomposition*.
- *Trie* is a data structure based on *key space decomposition*.

(a) Binary trie

(b) alphabet trie


The diagram illustrates two types of trie structures. (a) Binary trie: A binary tree where internal nodes are black dots. Edges are labeled 0 and 1. Leaf nodes are labeled with integers: 2, 7, 24, 32, 37, 40, 42, and 120. (b) Alphabet trie: A tree where internal nodes are black dots. Edges are labeled with characters. Leaf nodes are labeled with words: ant, anteater, antelope, chicken, deer, duck, goat, goldfish, and goose.

2. Balanced Trees

- **Balanced** may be defined by
 - *Height balancing*: comparing the heights of the two sub trees
 - *Null-path-length balancing*: comparing the null-path-length of each of the two sub-trees (the length to the closest null sub-tree/empty node)
 - *Weight balancing*: comparing the number of null sub-trees in each of the two sub trees
- Balance will ensure the height is $\Theta(\log n)$

Balanced trees - AVL tree


- AVL trees use height balancing
 - For every node, the heights of its left and right subtrees differ by **at most 1**.



AVL trees with the height of 4


Balanced trees - Red-black tree

- Red-black trees use null-path-length balancing
 - All nodes are colored red or black (0 or 1)
 - The root must be black
 - All children of a red node must be black
 - Any path from the root to an empty node must have the same number of black nodes
 - **Length**: One sub-tree must not be greater than twice the other.



Balanced trees - $\text{BB}(\alpha)$ tree


- $\text{BB}(\alpha)$ trees ($0 < \alpha \leq 1/3$) use weight balancing
 - Neither side has less than a proportion α of the empty nodes, i.e., both proportions fall in $[\alpha, 1 - \alpha]$



$\text{BB}(5/10)$ trees

Balanced trees - Splay tree

- Splay tree falls into an average cost $\Theta(\log n)$ of per access operation.
 - Access nodes could be rotated or *splayed* to the root of the tree.



Splay tree single rotation

3. Spatial data structures

- Searching on a **one-dimensional** key
 - BST, AVL tree, splay tree, 2-3 tree, B-tree, tries
- Searching on **multi-dimensional** key
 - Requires the use of spatial data structure
- **Spatial data structure** store data objects in two or more dimensions
 - widely used in geographic information systems, computer graphics, robotics, etc.

3. Spatial data structures (cont'd)


- Two typical Spatial Data Structures
 - K-D tree
 - PR quadtree

K-D tree


- K-D tree is **an extension of the BST** to multiple dimensions.
 - It is a binary tree whose splitting decisions alternate among the key dimensions.
 - Like the BST, the K-D tree uses **object space decomposition**.

K-D tree (cont'd)

- In a K-D tree, each level makes branching decisions based on a particular search key associated with that level.



(a)





(b)

PR quadtree

- PR (**P**oint-**R**egion) quadtree is a form of **trie**.
 - It uses **key space decomposition**.
 - It is a binary tree only for one-dimensional keys (in which case it is a trie with a binary alphabet).
 - For d dimensions it has 2^d branches. Thus, in two dimensions, the PR quadtree has four branches (hence the name "quadtree"), splitting space into four equal-sized quadrants at each branch.

PR quadtree (cont'd)

- In a PR quadtree, each node either has exactly four children or is a leaf.
- A full four-way branching (4-ary) tree in shape.



Summary

- Indexing
 - Tree-based index
 - B trees, B⁺ trees,
- Advanced Trees
 - Tries
 - Balanced trees
 - AVL tree, Red-black tree, BB(α) tree, Splay tree
 - Spatial data structures
 - K-D tree, PR quadtree